

Cookpad ML Researcher Technical Test

Imagine that we are building a service that allows our users to distinguish between different types of food using a smartphone camera.

Please implement a deep learning model that, given a photo, will classify it into one of two categories: *sushi* or *sandwich*. Please write your code in Python, but if not - please justify why your choice is better in this case. You are free to use Tensorflow, Keras, Chainer, Caffe or any other popular machine learning framework. You may adapt and modify an open-source project, as long as you respect its license.

In your project, please include a **Dockerfile** to build a docker image that will install all the dependencies, download data and run the training and validation scripts. It's also acceptable if the docker image runs a Jupyter notebook server.

In **README.md**:

- please analyze your model's accuracy, precision, recall or any other metrics you find interesting,
- write a short analysis on how the model can be used in production and what modifications would be necessary to deploy it into a useful product.

Since the problem is computationally intensive and time is short, your choice of framework, model and analysis of the results are more important than the final accuracy you achieve. We hope you make your choices with deployment in mind.

Key points:

- please make sure your code is legible and instructions how to run it are clear,
- you may give us the solution in a git repository that we can access.

Data that can be used for training or testing the model:

http://research.us-east-1.s3.amazonaws.com/public/sushi_or_sandwich_photos.zip

You are free to use data from other sources too, if you need.